



# Machine Learning

**Dr. Shahid Mahmood Awan**

**Assistant Professor**

**School of Systems and Technology, University of Management and Technology**

**shahid.awan@umt.edu.pk**

**Umer Saeed**(MS Data Science, BSc Telecommunication Engineering)

**Sr. RF Optimization & Planning Engineer**

**f2017313017@umt.edu.pk**



# Machine Learning

## Linear Regression



# Machine Learning

**Data Source**

# Energy Efficiency Data Set

---

- ▶ Energy Efficiency Data Set selected for Task-1 from the following URL;  
<https://archive.ics.uci.edu/ml/datasets/Energy+efficiency>
- ▶ **Source:**
- ▶ The dataset was created by Angeliki Xifara and was processed by Athanasios Tsanas.

# Energy Efficiency Data Set

---

## ► Data Set Information:

- We perform energy analysis using 12 different building shapes simulated in Ecotect.
- The buildings differ with respect to the glazing area, the glazing area distribution, and the orientation, amongst other parameters.
- 
- We simulate various settings as functions of the afore-mentioned characteristics to obtain 768 building shapes.
- The dataset comprises 768 samples and 8 features, aiming to predict two real valued responses. It can also be used as a multi-class classification problem if the response is rounded to the nearest integer.

# Energy Efficiency Data Set

## ► Attribute Information:

- The dataset contains eight attributes (or features, denoted by X1...X8) and two responses (or outcomes, denoted by Y1 and Y2). The aim is to use the eight features to predict each of the two responses.

Attributes/Features/Variables
<b>X1</b> Relative Compactness
<b>X2</b> Surface Area
<b>X3</b> Wall Area
<b>X4</b> Roof Area
<b>X5</b> Overall Height
<b>X6</b> Orientation
<b>X7</b> Glazing Area
<b>X8</b> Glazing Area Distribution

Outcome/Responses
<b>Y1</b> Heating Load
<b>Y2</b> Cooling Load

# Energy Efficiency Data Set

Item	Variable	Discrete Values	Range
1	X1 Relative Compactness	12	[0.62 , 0.98]
2	X2 Surface Area	12	[514.5 , 808.5]
3	X3 Wall Area	7	[245 , 416.5]
4	X4 Roof Area	4	[110.25 , 220.5]
5	X5 Overall Height	2	[3.5 , 7]
6	X6 Orientation	4	[2 ,5]
7	X7 Glazing Area	4	[0 ,0.4]
8	X8 Glazing Area Distribution	6	[0 , 5]

► **Table 1: Input Variables (Attributs)**

Item	Outcome	Mean	Range
1	Y1 Heating Load	22.30720052	[6.01 , 43.1]
2	Y2 Cooling Load	24.58776042	[10.9, 48.03]

► **Table 2: Output Variables (Responses)**



# Machine Learning

**Linear Regression with one variable**



# Linear Regression with One Variable

---

- ▶ The first step is to import the required Libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import scipy.stats as stats
```

- ▶ The second step is to load the dataset (**Energy Efficiency Data Set**)

```
enb=pd.read_csv("ENB2012_data.csv")
```


- ▶ The data will be loaded using Python Pandas, a data analysis module. It will be loaded into a structure known as a Panda Data Frame, which allows for each manipulation of the rows and columns.



# Machine Learning

**Linear Regression with One Variable(Outcome  
Heating Loading)**

# Linear Regression with One Variable



- Create two arrays: x (X1 Relative Compactness) and y (Y1 Heating Load ). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable

---

- ▶ In next step, Splitting the dataset into the features set and Outcome set

```
x = enb.iloc[:, 0].values
```

```
x=x.reshape(len(x),1)
```

```
y = enb.iloc[:, -2].values
```

```
y=y.reshape(len(y),1)
```

- ▶ The data will be split into a training and test set. Once we have the test data, we can find a best fit line and make predictions.

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.1, random_state= 0)
```

# Linear Regression with One Variable

---

- ▶ Fitting Simple Linear Regression to the Training set

```
from sklearn.linear_model import LinearRegression
```

```
regressor = LinearRegression()
```

```
regressor.fit(x_train, y_train)
```

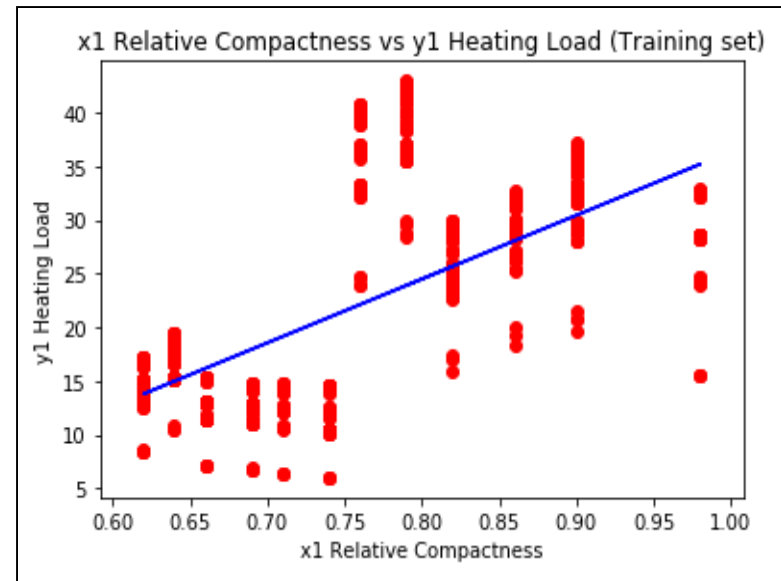
- ▶ Predicting the Test set results

```
y_pred = regressor.predict(x_test)
```

# Linear Regression with One Variable

- ▶ Visualising the Training set results

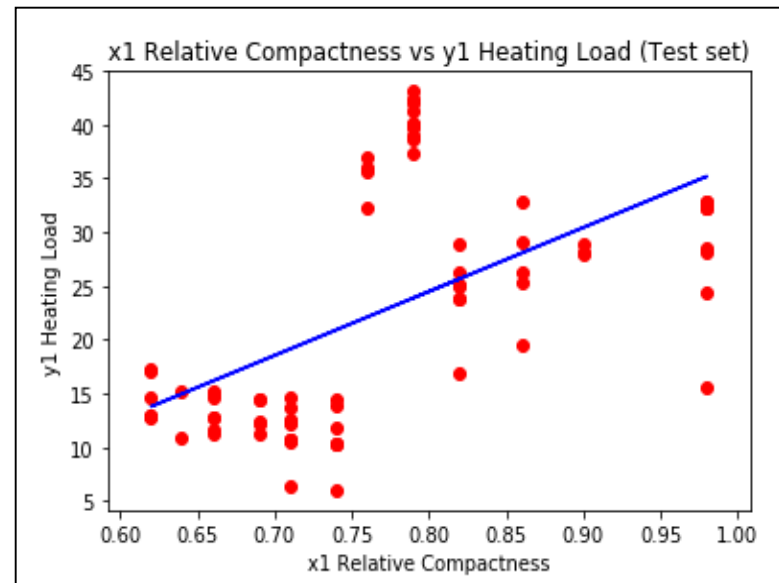
```
plt.scatter(x_train, y_train, color = 'red')  
plt.plot(x_train, regressor.predict(x_train), color = 'blue')  
plt.title('x1 Relative Compactness vs y1 Heating Load (Training set)')  
plt.xlabel('x1 Relative Compactness')  
plt.ylabel('y1 Heating Load')  
plt.show()
```



# Linear Regression with One Variable

- ▶ Visualising the Test set results

```
plt.scatter(x_test, y_test, color = 'red')  
plt.plot(x_train, regressor.predict(x_train), color = 'blue')  
plt.title('x1 Relative Compactness vs y1 Heating Load (Test set)')  
plt.xlabel('x1 Relative Compactness')  
plt.ylabel('y1 Heating Load')  
plt.show()
```



# Linear Regression with One Variable

---

- ▶ Building the optimal model using Backward Elimination

```
import statsmodels.formula.api as sm
```

```
x= np.append(arr=np.ones((768,1)).astype(int),values=x,axis=1)
```

```
x_opt = x[:,[0,1]]
```

```
regressor_OSL = sm.OLS(endog=y, exog=x_opt).fit()
```

```
regressor_OSL.summary()
```



# Linear Regression with One Variable

```

=====
OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.387
Model:                  OLS    Adj. R-squared:           0.386
Method:                 Least Squares    F-statistic:             484.0
Date:                  Sun, 22 Apr 2018    Prob (F-statistic):       1.59e-83
Time:                  12:35:23    Log-Likelihood:          -2676.5
No. Observations:      768    AIC:                     5357.
Df Residuals:          766    BIC:                     5366.
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-23.0530	2.081	-11.076	0.000	-27.139	-18.967
x1	59.3590	2.698	22.001	0.000	54.063	64.655

```


=====
Omnibus:                53.989    Durbin-Watson:           0.305
Prob(Omnibus):           0.000    Jarque-Bera (JB):        64.327
Skew:                    0.708    Prob(JB):                1.08e-14
Kurtosis:                3.059    Cond. No.                 15.0
=====

```

Warnings:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

# Linear Regression with One Variable




- Create two arrays: x (X2 Surface Area) and y (Y1 Heating Load ). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $X_3$  Wall Area) and  $y$  ( $Y_1$  Heating Load). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $X_4$  Roof Area) and  $y$  ( $Y_1$  Heating Load). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $X_5$  Overall Height) and  $y$  ( $Y_1$  Heating Load ). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $X_6$  Orientation) and  $y$  ( $Y_1$  Heating Load ). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression



- Create two arrays: x (X7 Glazing Area) and y (Y1 Heating Load ). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression



- Create two arrays:  $x$  ( $X_8$  Glazing Area Distribution) and  $y$  ( $Y_1$  Heating Load ). Intuitively we'd expect to find some correlation between the two variables.






# Machine Learning


**Linear Regression with One Variable(Outcome Cooling Loading)**

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $x_1$  Relative Compactness) and  $y$  ( $y_2$  Cooling Load ). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $X_2$  Surface Area) and  $y$  ( $Y_2$  Cooling Load ). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $X_3$  Wall Area) and  $y$  ( $Y_2$  Cooling Load). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $X_4$  Roof Area) and  $y$  ( $Y_2$  Cooling Load). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable




- Create two arrays:  $x$  ( $X_5$  Overall Height) and  $y$  ( $Y_2$  Cooling Load ). Intuitively we'd expect to find some correlation between the two variables.

# Linear Regression with One Variable



- Create two arrays: x (X6 Orientation) and y (Y2 Cooling Load ). Intuitively we'd expect to find some correlation between the two variables.


# Linear Regression



- Create two arrays: x (X7 Glazing Area) and y (Y2 Cooling Load ). Intuitively we'd expect to find some correlation between the two variables.



# Linear Regression



- Create two arrays:  $x$  ( $X_8$  Glazing Area Distribution) and  $y$  ( $Y_2$  Cooling Load ). Intuitively we'd expect to find some correlation between the two variables.



# Machine Learning

## Linear Regression with Multiple Variables

# Linear Regression with One Variable

---

- ▶ The first step is to import the required Libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import scipy.stats as stats
```

- ▶ The second step is to load the dataset (**Energy Efficiency Data Set**)

```
enb=pd.read_csv("ENB2012_data.csv")
```


- ▶ The data will be loaded using Python Pandas, a data analysis module. It will be loaded into a structure known as a Panda Data Frame, which allows for each manipulation of the rows and columns.



# Machine Learning

**Linear Regression with Multiple Variables  
(Outcome Heating Loading)**

# Linear Regression with Multiple Variables



$$h_{\theta}(x) = \theta_0 x_0 + \theta_{x1} X1 + \theta_{x2} X2 + \theta_{x3} X3 + \theta_{x4} X4 + \theta_{x5} X5 + \theta_{x6} X6 + \theta_{x7} X7 + \theta_{x8} X8$$

# Linear Regression with Multiple Variables

---



- ▶ In next step, Splitting the dataset into the features set and Outcome set

```
x = enb.iloc[:, :8].values
```

```
y = enb.iloc[:, 8].values
```

- ▶ The data will be split into a training and test set. Once we have the test data, we can find a best fit line and make predictions.

```
from sklearn.cross_validation import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.1, random_state= 0)
```

# Linear Regression with Multiple Variables



- ▶ Fitting Multiple Linear Regression to the Training set

```
from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(X_train, y_train)
```

- ▶ Predicting the Test set results

```
y_pred = regressor.predict(x_test)
```

- ▶ Building the optimal model using Backward Elimination

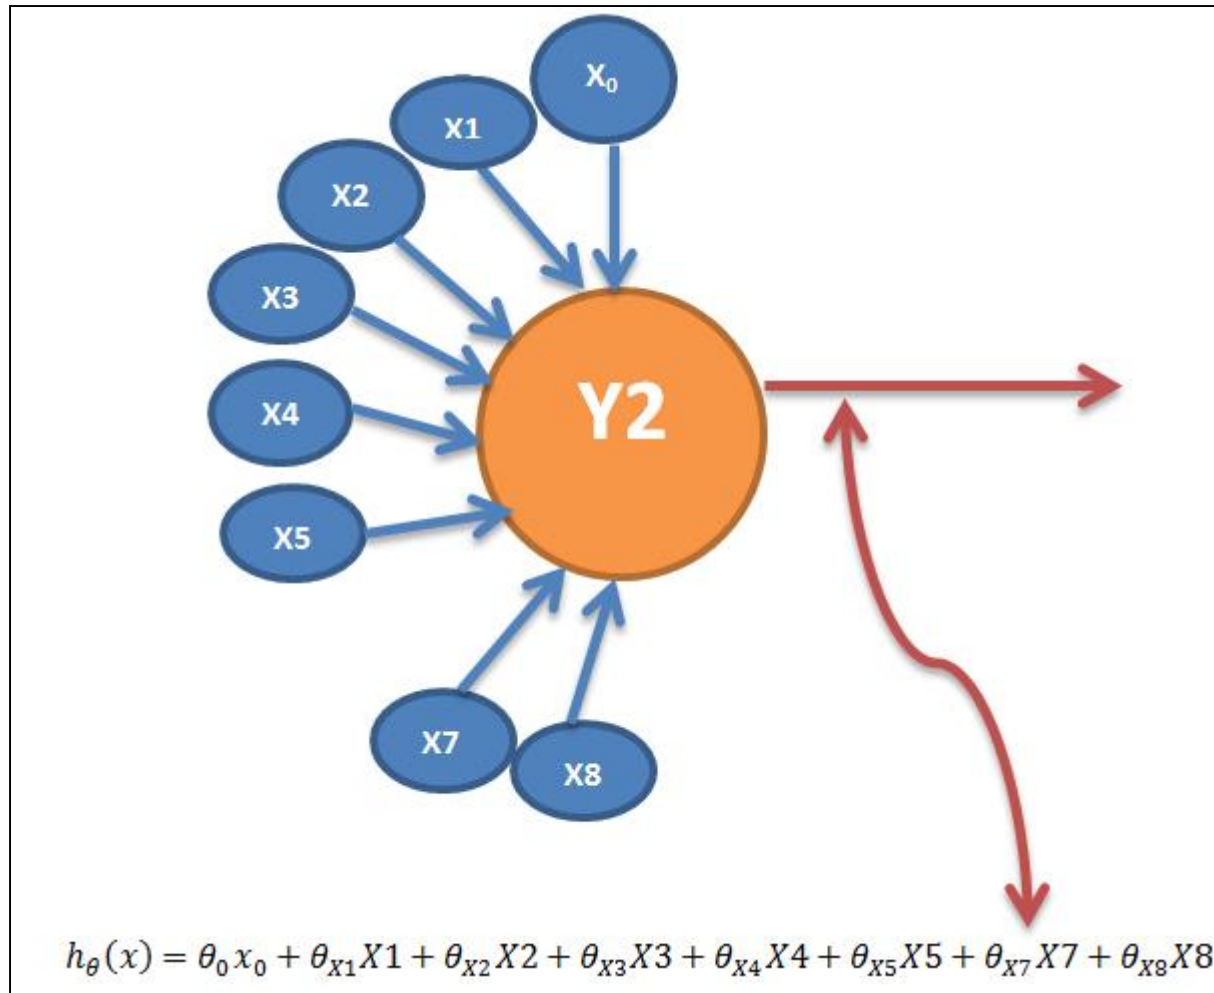
```
import statsmodels.formula.api as sm  
x= np.append(arr=np.ones((768,1)).astype(int),values=x,axis=1)  
x_opt = x[:,[0,1,2,3,4,5,6,7,8]]  
regressor_OLS = sm.OLS(endog=y, exog=x_opt).fit()  
regressor_OLS.summary()
```

# Linear Regression with Multiple Variables

OLS Regression Results						
=====						
Dep. Variable:	y	R-squared:	0.916			
Model:	OLS	Adj. R-squared:	0.915			
Method:	Least Squares	F-statistic:	1187.			
Date:	Sun, 22 Apr 2018	Prob (F-statistic):	0.00			
Time:	13:01:09	Log-Likelihood:	-1912.5			
No. Observations:	768	AIC:	3841.			
Df Residuals:	760	BIC:	3878.			
Df Model:	7					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	84.0145	19.034	4.414	0.000	46.650	121.379
x1	-64.7740	10.289	-6.295	0.000	-84.973	-44.575
x2	-0.0626	0.013	-4.670	0.000	-0.089	-0.036
x3	0.0361	0.004	9.386	0.000	0.029	0.044
x4	-0.0494	0.008	-6.569	0.000	-0.064	-0.035
x5	4.1699	0.338	12.337	0.000	3.506	4.833
x6	-0.0233	0.095	-0.246	0.805	-0.209	0.163
x7	19.9327	0.814	24.488	0.000	18.335	21.531
x8	0.2038	0.070	2.914	0.004	0.067	0.341
=====						
Omnibus:	18.648	Durbin-Watson:	0.654			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	37.708			
Skew:	0.044	Prob(JB):	6.48e-09			
Kurtosis:	4.082	Cond. No.	3.34e+15			
=====						
Warnings:						
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.						
[2] The smallest eigenvalue is 4.08e-23. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.						



# Linear Regression with Multiple Variables



# Linear Regression with Multiple Variables



- ▶ Building the optimal model using Backward Elimination

```
import statsmodels.formula.api as sm
x= np.append(arr=np.ones((768,1)).astype(int),values=x,axis=1)
x_opt = x[:,[0,1,2,3,4,5,7,8]]
regressor_OSL = sm.OLS(endog=y, exog=x_opt).fit()
regressor_OSL.summary()
```

# Linear Regression with Multiple Variables



```
=====
                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.916
Model:                  OLS    Adj. R-squared:           0.916
Method:                 Least Squares    F-statistic:       1387.
Date:                   Sun, 22 Apr 2018    Prob (F-statistic): 0.00
Time:                   13:03:59    Log-Likelihood:    -1912.5
No. Observations:      768    AIC:                3839.
Df Residuals:          761    BIC:                3871.
Df Model:               6
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                83.9329      19.019      4.413      0.000      46.597      121.269
x1                   -64.7740      10.283     -6.299      0.000     -84.961     -44.587
x2                   -0.0626      0.013     -4.673      0.000     -0.089     -0.036
x3                    0.0361      0.004      9.392      0.000      0.029      0.044
x4                   -0.0494      0.008     -6.573      0.000     -0.064     -0.035
x5                    4.1699      0.338     12.345      0.000      3.507      4.833
x6                   19.9327      0.813     24.503      0.000     18.336     21.530
x7                    0.2038      0.070      2.916      0.004      0.067      0.341
=====
Omnibus:               18.654    Durbin-Watson:           0.654
Prob(Omnibus):         0.000    Jarque-Bera (JB):        37.740
Skew:                  0.044    Prob(JB):                6.38e-09
Kurtosis:              4.082    Cond. No.                1.26e+16
=====


Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The smallest eigenvalue is 2.85e-24. This might indicate that there are
strong multicollinearity problems or that the design matrix is singular.
```



# Machine Learning

**Linear Regression with Multiple Variables (Outcome Cooling Loading)**

# Linear Regression with Multiple Variables



$$h_{\theta}(x) = \theta_0 x_0 + \theta_{x_1} X_1 + \theta_{x_2} X_2 + \theta_{x_3} X_3 + \theta_{x_4} X_4 + \theta_{x_5} X_5 + \theta_{x_6} X_6 + \theta_{x_7} X_7 + \theta_{x_8} X_8$$

# Linear Regression with Multiple Variables

---



- ▶ In next step, Splitting the dataset into the features set and Outcome set

```
x = enb.iloc[:, :8].values
```

```
y = enb.iloc[:, 9].values
```

- ▶ The data will be split into a training and test set. Once we have the test data, we can find a best fit line and make predictions.

```
from sklearn.cross_validation import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.1, random_state= 0)
```

# Linear Regression with Multiple Variables



- ▶ Fitting Multiple Linear Regression to the Training set

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

- ▶ Predicting the Test set results

```
y_pred = regressor.predict(x_test)
```

- ▶ Building the optimal model using Backward Elimination

```
import statsmodels.formula.api as sm
x= np.append(arr=np.ones((768,1)).astype(int),values=x,axis=1)
x_opt = x[:,[0,1,2,3,4,5,6,7,8]]
regressor_OLS = sm.OLS(endog=y, exog=x_opt).fit()
regressor_OLS.summary()
```

# Linear Regression with Multiple Variables

```

=====
                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.888
Model:                  OLS    Adj. R-squared:            0.887
Method:                 Least Squares    F-statistic:        859.1
Date:                   Sun, 22 Apr 2018    Prob (F-statistic):    0.00
Time:                   13:21:08    Log-Likelihood:       -1979.3
No. Observations:       768    AIC:                  3975.
Df Residuals:           760    BIC:                  4012.
Df Model:                7
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	97.2457	20.765	4.683	0.000	56.483	138.009
x1	-70.7877	11.225	-6.306	0.000	-92.824	-48.751
x2	-0.0661	0.015	-4.519	0.000	-0.095	-0.037
x3	0.0225	0.004	5.365	0.000	0.014	0.031
x4	-0.0443	0.008	-5.404	0.000	-0.060	-0.028
x5	4.2838	0.369	11.618	0.000	3.560	5.008
x6	0.1215	0.103	1.176	0.240	-0.081	0.324
x7	14.7171	0.888	16.573	0.000	12.974	16.460
x8	0.0407	0.076	0.534	0.594	-0.109	0.190

```

=====
Omnibus:                 104.668    Durbin-Watson:           1.094
Prob(Omnibus):           0.000    Jarque-Bera (JB):        230.547
Skew:                    0.767    Prob(JB):                 8.65e-51
Kurtosis:                 5.203    Cond. No.                 3.34e+15
=====

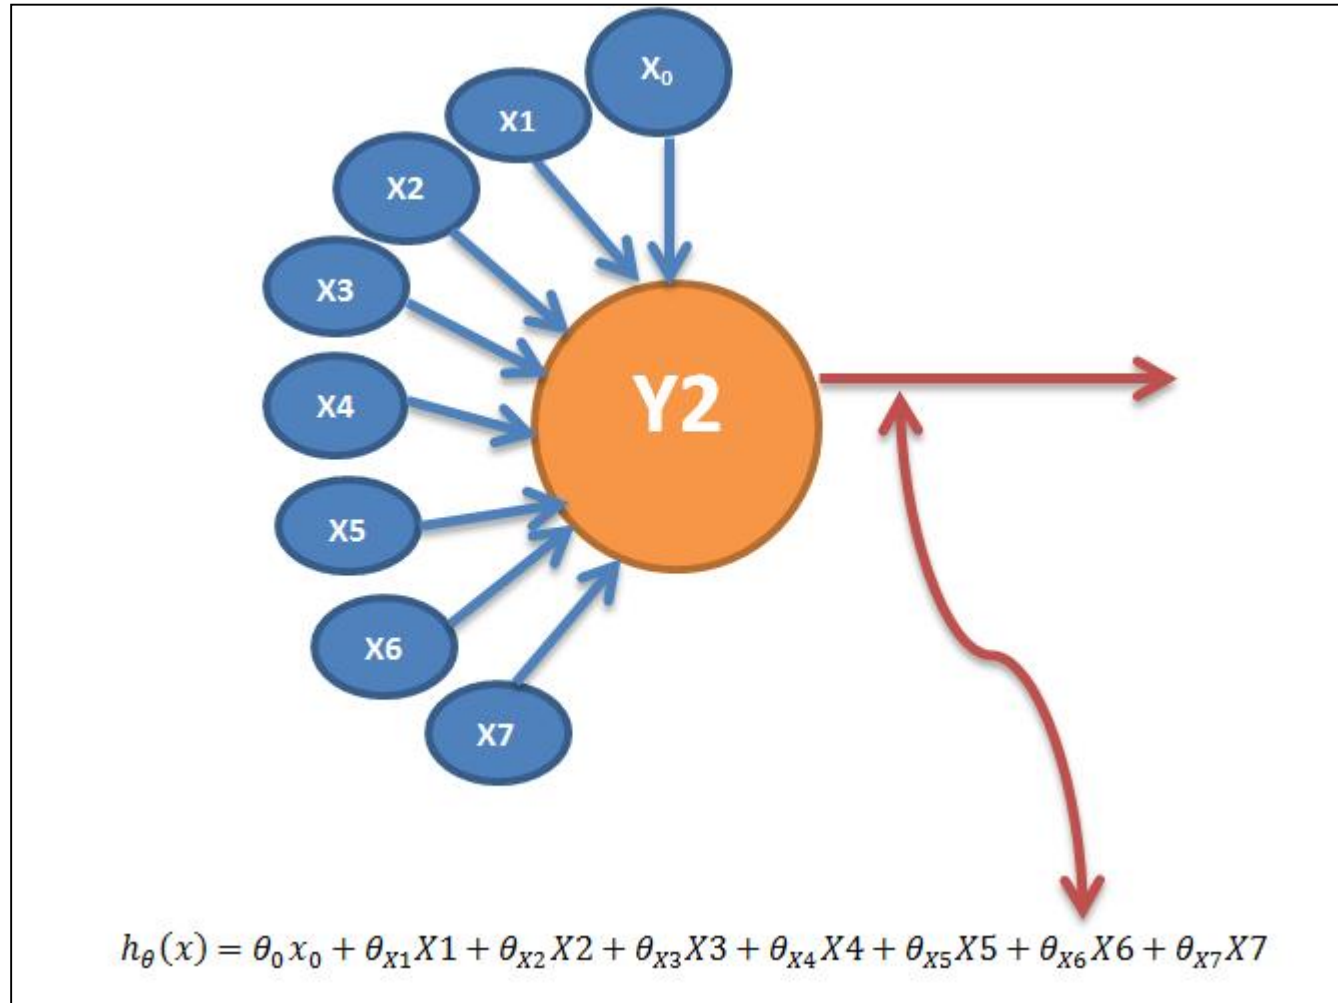
```

## Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 4.08e-23. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.



# Linear Regression with Multiple Variables



# Linear Regression with Multiple Variables




- ▶ Building the optimal model using Backward Elimination

```
import statsmodels.formula.api as sm
x= np.append(arr=np.ones((768,1)).astype(int),values=x,axis=1)
x_opt = x[:,[0,1,2,3,4,5,6,7]]
regressor_OSL = sm.OLS(endog=y, exog=x_opt).fit()
regressor_OSL.summary()
```

# Linear Regression with Multiple Variables

OLS Regression Results						
=====						
Dep. Variable:	y	R-squared:	0.888			
Model:	OLS	Adj. R-squared:	0.887			
Method:	Least Squares	F-statistic:	1003.			
Date:	Sun, 22 Apr 2018	Prob (F-statistic):	0.00			
Time:	13:29:24	Log-Likelihood:	-1979.5			
No. Observations:	768	AIC:	3973.			
Df Residuals:	761	BIC:	4005.			
Df Model:	6					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	97.3366	20.754	4.690	0.000	56.594	138.079
x1	-70.7877	11.220	-6.309	0.000	-92.814	-48.762
x2	-0.0661	0.015	-4.521	0.000	-0.095	-0.037
x3	0.0225	0.004	5.367	0.000	0.014	0.031
x4	-0.0443	0.008	-5.407	0.000	-0.060	-0.028
x5	4.2838	0.369	11.623	0.000	3.560	5.007
x6	0.1215	0.103	1.177	0.240	-0.081	0.324
x7	14.8180	0.867	17.086	0.000	13.116	16.520
=====						
Omnibus:	104.448	Durbin-Watson:	1.093			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	230.721			
Skew:	0.764	Prob(JB):	7.93e-51			
Kurtosis:	5.208	Cond. No.	1.26e+16			
=====						
Warnings:						
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.						
[2] The smallest eigenvalue is 2.85e-24. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.						

# Linear Regression with Multiple Variables



$$h_{\theta}(x) = \theta_0 x_0 + \theta_{x_1} X_1 + \theta_{x_2} X_2 + \theta_{x_3} X_3 + \theta_{x_4} X_4 + \theta_{x_5} X_5 + \theta_{x_7} X_7$$

# Linear Regression with Multiple Variables

---

- ▶ Building the optimal model using Backward Elimination

```
import statsmodels.formula.api as sm
x= np.append(arr=np.ones((768,1)).astype(int),values=x,axis=1)
x_opt = x[:,[0,1,2,3,4,5,7]]
regressor_OSL = sm.OLS(endog=y, exog=x_opt).fit()
regressor_OSL.summary()
```

# Linear Regression with Multiple Variables

OLS Regression Results						
=====						
Dep. Variable:	y	R-squared:	0.888			
Model:	OLS	Adj. R-squared:	0.887			
Method:	Least Squares	F-statistic:	1203.			
Date:	Sun, 22 Apr 2018	Prob (F-statistic):	0.00			
Time:	13:33:40	Log-Likelihood:	-1980.2			
No. Observations:	768	AIC:	3972.			
Df Residuals:	762	BIC:	4000.			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	97.7618	20.756	4.710	0.000	57.015	138.508
x1	-70.7877	11.223	-6.307	0.000	-92.819	-48.756
x2	-0.0661	0.015	-4.520	0.000	-0.095	-0.037
x3	0.0225	0.004	5.366	0.000	0.014	0.031
x4	-0.0443	0.008	-5.405	0.000	-0.060	-0.028
x5	4.2838	0.369	11.620	0.000	3.560	5.008
x6	14.8180	0.867	17.082	0.000	13.115	16.521
=====						
Omnibus:	104.896	Durbin-Watson:	1.095			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	232.225			
Skew:	0.766	Prob(JB):	3.74e-51			
Kurtosis:	5.215	Cond. No.	1.26e+16			
=====						
Warnings:						
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.						
[2] The smallest eigenvalue is 2.85e-24. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.						